

```
[2]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
# Step 1: Load the dataset
# Using the Titanic dataset as an example
data = pd.read_csv('https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv')
# Display the first few rows of the dataset
print("Dataset preview:")
print(data.head())
```

Dataset preview:

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```

# Step 2: Preprocess the data
# Selecting relevant features and target
selected_features = ['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare']
data = data[selected_features + ['Survived']]
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1}) # Encode categorical variable
data = data.dropna() # Remove rows with missing values
X = data[selected_features] # Features
y = data['Survived'] # Target
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Initialize and train the Logistic Regression model
model = LogisticRegression()
model.fit(X_train, y_train)

```

▼ LogisticRegression

LogisticRegression()

```
[12]: # Step 4: Make predictions
y_pred = model.predict(X_test)
# Display predictions for the first 10 test samples
print("\nSample predictions:")
print(f"Predicted: {y_pred[:10]}")
print(f"Actual:    {y_test.values[:10]}")
```

Sample predictions:

Predicted: [0 0 1 1 0 0 0 0 1 1]

Actual: [0 1 1 1 0 1 1 1 0 0]

```
[14]: # Step 5: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"\nAccuracy on test data: {accuracy * 100:.2f}%")
# Additional performance metrics
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy on test data: 74.83%

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.82	0.80	87
1	0.69	0.64	0.67	56
accuracy			0.75	143
macro avg	0.74	0.73	0.73	143
weighted avg	0.75	0.75	0.75	143

```
[20]: # Visualizations
# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Not Survived', 'Survived'], yticklabels=['Not Survived', 'Survived'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Distribution of predictions
sns.countplot(x=y_pred)
plt.title('Distribution of Predictions')
plt.xlabel('Predicted Class')
plt.ylabel('Count')
plt.show()
```

